

Viral Pneumonia Screening on Chest X-ray Images Using Confidence-Aware Anomaly Detection

Jianpeng Zhang, Yutong Xie, Zhibin Liao, Guansong Pang, Johan Verjans, Wenxin Li, Zongji Sun, Jian He, Yi Li, Chunhua Shen, and Yong Xia

Abstract—Cluster of viral pneumonia occurrences during a short period of time may be a harbinger of an outbreak or pandemic, like SARS, MERS, and recent COVID-19. Rapid and accurate detection of viral pneumonia using chest X-ray can be significantly useful in large-scale screening and epidemic prevention, particularly when other chest imaging modalities are less available. Viral pneumonia often have diverse causes and exhibit notably different visual appearances on X-ray images. The evolution of viruses and the emergence of novel mutated viruses further result in substantial dataset shift, which greatly limits the performance of classification approaches. In this paper, we formulate the task of differentiating viral pneumonia from non-viral pneumonia and healthy controls into an one-class classification-based anomaly detection problem, and thus propose the confidence-aware anomaly detection (CAAD) model, which consists of a shared feature extractor, an anomaly detection module, and a confidence prediction module. If the anomaly score produced by the anomaly detection module is large enough or the confidence score estimated by the confidence prediction module is small enough, we accept the input as an anomaly case (*i.e.*, viral pneumonia). The major advantage of our approach over binary classification is that we avoid modeling individual viral pneumonia classes explicitly and treat all known viral pneumonia cases as anomalies to reinforce the one-class model. The proposed model outperforms binary classification models on the clinical X-VIRAL dataset that contains 5,977 viral pneumonia (no COVID-19) cases, 18,619 non-viral pneumonia cases, and 18,774 healthy controls. Moreover, when directly testing on the X-COVID dataset that contains 106 confirmed COVID-19 cases and 107 normal controls without any fine-tuning, our model achieves an AUC of 83.61% and a sensitivity of 71.70%, comparable to the performance of radiologists reported in the literature.

Index Terms—Viral pneumonia screening; deep anomaly detection; confidence prediction; chest X-ray.

I. INTRODUCTION

This work was supported in part by the National Natural Science Foundation of China under Grants 61771397 and in part by the Science and Technology Innovation Committee of Shenzhen Municipality, China, under Grants JCYJ20180306171334997. (*Corresponding authors:* C. Shen and Y. Xia)

J. Zhang, Y. Xie, and Y. Xia are with the National Engineering Laboratory for Integrated Aero-Space-Ground-Ocean Big Data Application Technology, School of Computer Science and Engineering, Northwestern Polytechnical University, Xi'an 710072, China. Z. Liao, G. Pang, J. Verjans, and C. Shen are with the University of Adelaide, Australia. Y. Li is with the GreyBird Ventures, LLC. W. Li, and Z. Sun are with the JF Healthcare Inc. J. He is with the Department of Radiology, Nanjing Drum Tower Hospital The Affiliated Hospital of Nanjing University Medical School. The authors declare that there is no conflict of interests regarding the publication of this article. The first two authors' contribution was made when visiting The University of Adelaide. Z. Liao, G. Pang, J. Verjans, C. Shen and their employer received no financial support for the research, authorship, and/or publication of this article. (e-mail: james.zhang@mail.nwpu.edu.cn; chunhua.shen@adelaide.edu.au; yxia@nwpu.edu.cn)

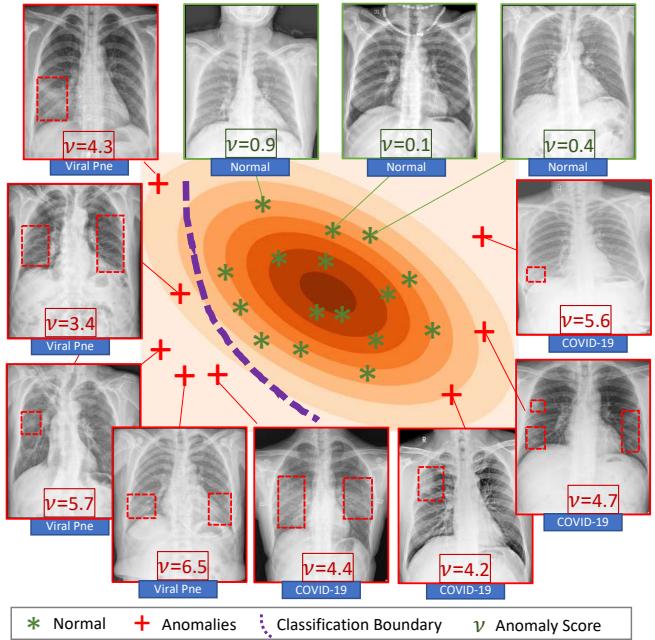


Fig. 1. An illustration of binary classification vs. anomaly detection in distinguishing viral pneumonia cases (*i.e.*, ‘Anomalies’) from non-viral cases and healthy controls (*i.e.*, ‘Normal’). Image samples are shown on a two-dimensional contour plot, where inner contour lines indicate large density. The red dotted boxes in viral pneumonia cases are the suspected lesion areas annotated by radiologists. Novel viral pneumonia (*e.g.*, COVID-19) can be either similar to or different from the cases caused by known viruses (denoted as ‘Viral Pne’). The decision boundary given by a binary classification approach can well separate ‘Normal’ cases and ‘Viral Pne’ cases, but may not able to distinguish ‘Normal’ cases from COVID-19 cases. In contrast, our anomaly detection approach is able to distinguish both known and novel viral pneumonia from ‘Normal’ cases by assigning the former large anomaly scores and the latter small anomaly scores.

IN contrast to bacterial pneumonia, viral pneumonia is a type of lung infections caused by viruses. These viruses invade lungs, causing them to swell and to block the flow of oxygen, which may be life-threatening. As a recent typical example, COVID-19, a viral pneumonia caused by a novel coronavirus, has rapidly spread worldwide within a few months, subsequently threatening the health of billions of human beings [1], [2]. The clustering occurrence of viral pneumonia, like SARS [3], MERS [4], and COVID-19, can often suggest a potential outbreak. Therefore, it is desirable to develop an accurate, fast, and cost-effective tool for viral pneumonia detection, which provides the prerequisite for rigorous detection, contact tracing, and isolation of infected subjects in a large district.

Taking the recent COVID-19 outbreak for example, viral nucleic acid detection using real-time polymerase chain reaction (RT-PCR) is the accepted standard diagnostic method [5], [6]. However, many hyper-endemic regions and countries are not able to provide sufficient RT-PCR testing for tens of thousands of suspected subjects in a short period of time. Moreover, it would fail to detect the newly evolved coronavirus before extracting the DNA sequence of the new virus, which may delay the control of the epidemic. Hence, medical imaging is considered as an important technique to assist doctors to evaluate the disease development and to make prevention and control measures as soon as possible. Clinically, chest X-ray is the most commonly used imaging modality in the diagnostic workup of patients with thoracic abnormalities, due to its fast imaging speed and significantly low cost [7]. In comparison to computed tomography (CT), chest X-ray cannot provide the 3D anatomy of human body, but is generally considered to be able to differentiate between viral and non-viral pneumonia. Nevertheless, chest X-ray is probably the hardest plain film to interpret correctly [8]. Accurate interpretation can greatly influence patient management in the acute setting, such as following the clustering occurrence of COVID-19. Therefore, we aim to develop an automated and accurate viral pneumonia screening technique using chest X-ray as a stopgap for early warning of clusters of an outbreak caused by COVID-19 or a novel coronavirus.

The recent advances in deep learning have led to breakthroughs in many long-standing medical image analysis tasks, such as the detection, staging, and delineation of pathological abnormalities. On the task of chest X-ray interpretation, deep convolutional neural networks (DCNNs) have been constructed to diagnose the most common and important thoracic diseases [9], [10], [11] and to differentiate between viral and bacterial pneumonia [12], [13]. In contrast, we attempt to distinguish the viral pneumonia from all non-viral pneumonia ones (not merely bacterial pneumonia), aiming to rapidly detect the clusters of viral pneumonia (e.g., COVID-19) caused by a novel virus before an outbreak. This task, however, remains challenging due to two intrinsic complexities. First, many types of viruses can cause pneumonia, including the influenza A/B viruses, respiratory syncytial virus, coronaviruses, herpes simplex, measles, chickenpox, and more seriously, some uncertain novel virus. The complex pathological cues of viral pneumonia usually cause considerable visual differences on X-ray images (see Fig. 1), leading to substantial intra-class variance and dataset shift (e.g., newly emerged viral pneumonia cases have highly different lesions from the known viral pneumonia in the training data). Second, it is hard to collect a large number of positive (*i.e.*, viral pneumonia) samples in the early stage of an outbreak. Hence, the viral versus non-viral pneumonia classification is an extreme class-imbalance problem. These two complexities pose significant challenges to the commonly-used binary classification approaches since (i) they normally assume that the training and test data share an identical distribution (*i.e.*, no dataset shift) and there exists small class variance within each class; and (ii) they often ignore the class-imbalance problem. Consequently, the classifiers yield poor sensitivity performance. The sensitivity,

however, is clinically significant, since it can be much more disastrous to discharge a patient with viral pneumonia than to misdiagnose a healthy control. To address both challenges, we advocate the replacement of a classifier by an anomaly detector for this chest X-ray interpretation problem. As an one-classification approach [14], anomaly detection is not only able to detect dissimilar or even previously unseen anomalies, but also less dependent on labeled anomaly data than binary classification [15], [16], [17].

In this paper, we propose a confidence-aware anomaly detection (CAAD) model to distinguish viral pneumonia cases from non-viral pneumonia cases and healthy controls using chest X-rays. We reformulate the viral pneumonia screening into a one-class classification-based anomaly detection task, instead of a binary classification one. Specifically, we design an anomaly detection module to assign each X-ray image an anomaly score and employ the contrastive loss function to ensure that the scores generated for anomalies (*i.e.*, viral pneumonia) are significantly larger than those for non-viral pneumonia cases and normal controls. We further introduce an additional confidence prediction module to describe the confidence of the anomaly detection module. Judging by the confidence level, we re-assign the samples with low confidence as suspected viral pneumonia for further medical tests, which is helpful in reducing false negative cases and thus improving the sensitivity. Both the anomaly detection module and confidence prediction module can be jointly optimized in an end-to-end manner. We have evaluated our CAAD model on the X-VIRAL dataset which contains 5,977 positive viral pneumonia subjects and 37,393 negative subjects. Our proposal achieves the state-of-the-art performance, *i.e.*, 87.57% AUC, for viral pneumonia screening. Even with no exposure to COVID-19 cases during training, our CAAD model shows superior performance for the COVID-19 screening purpose, achieving an AUC of 83.61% and a sensitivity of 71.70% on our additional unseen X-COVID dataset with 106 confirmed and 107 normal subjects. This is comparable to the performance of radiologists reported in the literature [18].

The main contributions of this paper are summarized as the followings: (1) we formulate the viral pneumonia screening into an anomaly detection problem and propose the CAAD model to solve it, which is able to detect viral pneumonia caused by novel viruses and is less dependent on labeled viral pneumonia data than classification models; (2) we propose to predict failures of anomaly detection by modeling its confidence level to further improve the screening sensitivity; and (3) our experimental results demonstrate the effectiveness and strong generalizability of our model in viral pneumonia screening and the potential on the epidemic prevention and control.

II. RELATED WORK

A. Chest X-ray for pulmonary disease screening

Chest X-ray is one of commonly used imaging modalities for visualizing and quantifying the structural and functional consequences of thoracic diseases, providing high-resolution pictures of disease progression and therapy response. Magree

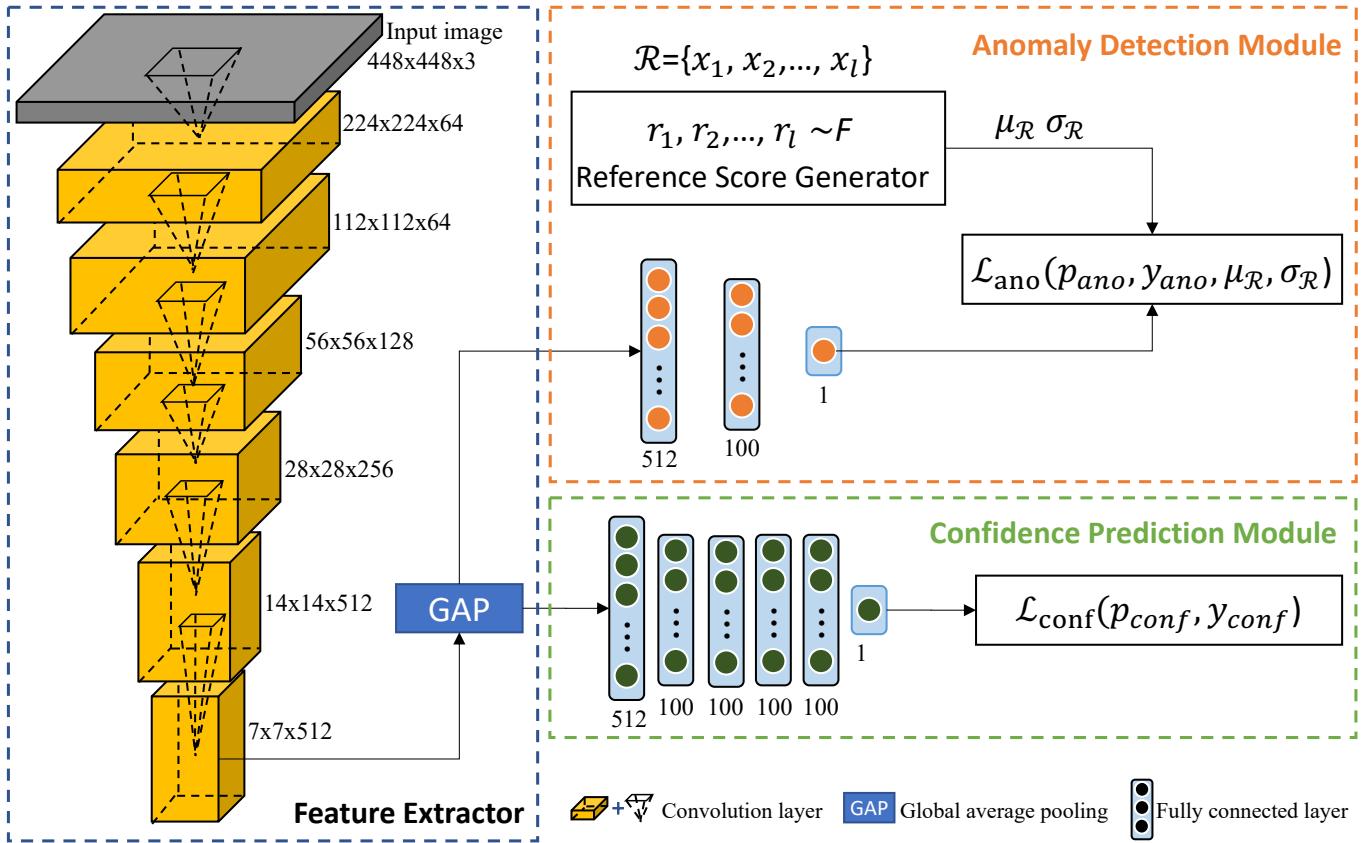


Fig. 2. Diagram of the proposed CAAD model. This model is composed of an anomaly detection module and a confidence prediction module, which are assigned to predict the anomaly score and confidence score of each input, respectively.

et al. [19] documented the incidence of pneumonia confirmed with X-ray imaging and demonstrated a high incidence, which guided the later prevention and treatment of vaccine. Jacobi *et al.* [20] described the most common manifestations and patterns of lung abnormality on chest X-ray in COVID-19 and suggested that medical community can frequently rely on portable chest X-ray due to its widespread availability and reduced infection control issues that currently limit CT utilization. Wong *et al.* [18] demonstrated that the common CT findings of bilateral involvement, peripheral distribution, and lower zone dominance can also be appreciated on chest X-ray, which shows the potential of using chest X-ray as a tool for identifying COVID-19. Borghesi *et al.* [21] presented an experimental chest X-ray scoring system and applied it to hospitalized patients with COVID-19 to quantify and monitor the severity and progression of COVID-19. Different from these studies, we focus on the viral pneumonia screening and aim to develop a fast and accurate algorithm to differentiate viral pneumonia from non-viral pneumonia and normal controls for the prevention and control of a possible outbreak.

B. Deep learning for chest X-ray interpretation

To improve efficiency and ease the burden of radiologists, researchers gradually adapt the recent advances of deep learning to interpret chest X-ray images. For the computer-aided diagnosis of 14 common thoracic diseases, Wang *et al.* [9] proposed a weakly-supervised classification and localization

framework, Rajpurkar *et al.* [10] constructed a 121-layer dense convolutional neural network that can perform the task at a level exceeding practicing radiologists, and Wang *et al.* [11] introduced an attention mechanism to help the model focus on the lesion area and thus further improved the diagnosis performance. Besides, many attempts [12], [13] have been made to investigate DCNN-based classification models for pneumonia detection and the differentiation between viral and bacterial pneumonia, aiming to facilitate rapid referrals for children who need urgent intervention. In these studies, these diagnostic tasks are formulated as classification problems, which are usually solved based on the intra-class similarity and inter-class dissimilarity of pathological patterns. Such classification models may fail to distinguish viral and non-viral pneumonia, since the category of viral pneumonia contains cases with highly variable visual appearances.

C. Failure prediction

Despite of their success, deep learning models still make mistakes, particularly when applied to real-world applications. To avoid the decision risk caused by the inherent defects of deep learning models, failure prediction is of great necessity. Hendrycks *et al.* [22] proposed to detect failures and out-of-distribution examples in neural networks via the prediction/maximum class probability method. However, it is hard to distinguish the failures if they are misclassified with a high probability. To address this issue, Corbiere *et al.* [23]

TABLE I

ARCHITECTURE OF OUR FEATURE EXTRACTOR. S: STAGE, L: NUMBER OF STACKED LAYERS, $W/H/C$: WEIGHT/HEIGHT/CHANNEL, GAP: GLOBAL AVERAGE POOLING

S	Operator	Input → Output ($W \times H \times C$)	L
1	Conv3x3	$448 \times 448 \times 3 \rightarrow 224 \times 224 \times 32$	1
	MBCConv1, k3x3	$224 \times 224 \times 32 \rightarrow 224 \times 224 \times 16$	1
2	MBCConv6, k3x3	$224 \times 224 \times 16 \rightarrow 112 \times 112 \times 24$	2
3	MBCConv6, k5x5	$112 \times 112 \times 24 \rightarrow 56 \times 56 \times 40$	2
	MBCConv6, k3x3	$56 \times 56 \times 40 \rightarrow 28 \times 28 \times 80$	3
4	MBCConv6, k5x5	$28 \times 28 \times 80 \rightarrow 28 \times 28 \times 112$	3
	MBCConv6, k5x5	$28 \times 28 \times 112 \rightarrow 14 \times 14 \times 192$	4
5	MBCConv6, k3x3	$14 \times 14 \times 192 \rightarrow 14 \times 14 \times 320$	1
	Conv3x3	$14 \times 14 \times 320 \rightarrow 7 \times 7 \times 320$	1
6	GAP	$7 \times 7 \times 320 \rightarrow 1 \times 1 \times 320$	1

considered the true class probability as a suitable confidence criterion for failure estimation. Xie *et al.* [24] proposed a deep segmentation-emendation model for gland instance segmentation, in which an emendation network is designed to predict the inconsistency between the ground truth masks and pixel-wise predictions of segmentation network, and the failure predictions made by an emendation network can be utilized to refine the segmentation result. Inspired by these works, we attempt to predict failures of anomaly detection, which improves not only the detection performance but also the diagnosis credibility of our model.

III. METHODS

The proposed CCAD model is composed of an anomaly detection network and a confidence prediction network (see Fig. 2). Both of networks share a feature extractor. Given an input chest X-ray image \mathbf{x} , the anomaly detection network aims to learn an anomaly scoring function $\varphi : \mathbf{x} \rightarrow \mathbb{R}$. For any two inputs \mathbf{x}_i and \mathbf{x}_j , we have $\varphi(\mathbf{x}_i) > \varphi(\mathbf{x}_j)$ if \mathbf{x}_i is abnormal and \mathbf{x}_j is normal. The confidence prediction network targets at approximating a confidence scoring function $\zeta : \mathbf{x} \rightarrow [0, 1]$, where 1 indicates highest model confidence and 0 indicates the opposite. In the inference mode, if the anomaly score is larger than T_{ano} or the confidence score is less than T_{conf} , we accept the input as an anomaly case (*i.e.*, viral pneumonia). We now delve into each part of our model.

A. Feature extractor

Although a DCNN with any architecture can be embedded in our CAAD model as the feature extractor, we choose the state-of-the-art EfficientNet [25] with the B0 architecture pretrained on ImageNet [26], due to the trade-off between the performance and complexity. This network is mainly composed of mobile inverted bottleneck (MBCConv) blocks [27], [28] with squeeze-and-excitation module [29]. The six stages of layer-by-layer convolution operations are represented by yellow rectangles in Figure 2, and the architecture details are listed in Table I. For each input chest X-ray image \mathbf{x} , it is first processed by several MBCConv blocks, and then transformed into a d -dimensional (d equals the number of channels in the last convolution layer) feature vector by a global average pooling layer. We denote the parameters of feature extractor as θ .

B. Anomaly detection network

The anomaly detection network is composed of the feature extractor and anomaly detection module that is a multi-layer perceptron with three 100-neuron hidden layers and an one-neuron output layer. It aims to generate an anomaly score for each input image \mathbf{x} , formulated as

$$\nu = \varphi(\mathbf{x}; \theta, \alpha) \quad (1)$$

where α is the trainable parameters of the anomaly detection module.

To guide the learning procedure of the anomaly detection module, we compute another scalar score as a reference. We randomly sample l normal data from a Gaussian distribution, *i.e.*, $r_1, r_2, \dots, r_l \sim \mathcal{N}(\mu, \sigma^2)$, and define a reference score as $\mu_R = \frac{1}{l} \sum_{i=1}^l r_i$. Following [17], we set $\mu = 0$ and $\sigma = 1$. With the obtained anomaly score and reference score, we employ the following contrastive loss [30], [31] to optimize the anomaly detection module

$$\begin{aligned} \mathcal{L}_{ano}(\nu, y, \mu_R, \sigma_R) = & (1 - y) \left| \frac{\nu - \mu_R}{\sigma_R} \right| \\ & + y \max \left(0, \text{margin} - \frac{\nu - \mu_R}{\sigma_R} \right) \end{aligned} \quad (2)$$

where σ_R is the standard deviation of the anomaly scores of randomly selected l normal data, y is the ground truth label, *i.e.*, $y = 0$ indicates that the input is a negative case and $y = 1$ indicates that the input is a positive case. Besides, margin represents the Z-score confidence interval parameter, which is empirically set to 5 for this study.

Different from binary classification whose performance can be largely degraded when there is imbalanced class distribution, our one-class classification-based anomaly detection network is inherently resilient to the class imbalance. This is because the anomaly detection network aims at learning a one-class description model from the large-scale negative data (*i.e.*, non-viral pneumonia). By doing so, it avoids modeling the positive class with the limited amount of labeled data; the limited positive data is used instead to reinforce the one-class modeling to achieve tighter one-class description.

C. Confidence prediction network

The current approach of anomaly detection does not have an error correction mechanism. However, we observe that the model do produce false predictions. To alleviate this issue, we follow the work of failure prediction in image classification and segmentation [23], [24], and make the shift in thinking that we can predict the failures of anomaly detection. Hence, we propose a confidence prediction network to learn a confidence score for each input, which reflects the confidence of anomaly score estimated by our model.

1) *Confidence criterion for anomaly detection:* The predicted anomaly score $\nu \in \mathbb{R}$ explicitly describes the abnormality degree of a given image, varying from very confirmed positive cases, *i.e.*, viral pneumonia ($\nu \geq \text{margin}$), to confirmed negative cases, *i.e.*, non-viral pneumonia or healthy patients ($\nu \approx 0$). However, it is difficult to describe the abnormality degree in the form of probability, which is important

for confidence prediction. Hence, we propose to employ the probability density function (PDF) to estimate the prediction probability. Given the former assumption that ν is drawn from $\mathcal{N}(\mu, \sigma^2)$, the Gaussian PDF is then:

$$\text{PDF}(\nu) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(\nu - \mu)^2}{2\sigma^2}\right), \quad (3)$$

We then approximate the prediction probability of anomaly detection by the normalized PDF, where PDF is commonly regarded as relative probability and the normalization scales the values into [0, 1]. The approximated prediction probability is expressed as:

$$\begin{aligned} \text{prob} &= \frac{\text{PDF}(\nu)}{\max(\text{PDF})}, \\ &= \exp\left(-\frac{(\nu - \mu)^2}{2\sigma^2}\right). \end{aligned} \quad (4)$$

However, it is hard to use such a prediction probability to distinguish failure predictions from successful ones (discussed in Section V-B). To further address this issue, we propose the anomaly probability, formulated as:

$$g = \begin{cases} \text{prob} & \text{if } \mathbf{x} \text{ is negative} \\ 1 - \text{prob} & \text{if } \mathbf{x} \text{ is positive} \end{cases}. \quad (5)$$

where $g \in [0, 1]$ is the confidence criterion for distinguishing successful and erroneous predictions. Intuitively, a robust model should successfully predict the true labels of input cases with a high confidence, whereas should have low confidence when making erroneous predictions. Therefore, in Eq. 5, g is close to 0 when the anomaly detection module fails to predict the true labels of given image and close to 1 when the true labels are successfully predicted.

2) *Confidence prediction network*: Our confidence prediction network is also built upon the shared feature extractor, and particularly contains the confidence prediction module with four 100-neuron hidden layers (see Fig. 2) for a strong prediction ability, as done in [23]. The forward computation of the confidence prediction network can be formally expressed as

$$\iota = \zeta(\mathbf{x}; \boldsymbol{\theta}, \boldsymbol{\beta}) \quad (6)$$

where ι is the confidence score of the corresponding anomaly detection, and $\boldsymbol{\beta}$ represents the ensemble of parameters of this module.

Since the confidence score ι takes a value from the range [0, 1], we formulate confidence prediction as a regression task and employ the standard l_2 loss to optimize the confidence prediction module.

$$\mathcal{L}_{conf}(\iota, g) = |\iota - g|^2 \quad (7)$$

D. Training and inference

We resize each training image to a fixed size of 512×512 pixels and applied several data augmentation strategies, including random cropping patches of size 448×448 and zooming (90%~110%) and horizontally flipping cropped patches, to alleviate overfitting on the training set. Then, the proposed CAAD model is trained in three steps. First, we train the

Algorithm 1: Training CAAD model

Input: $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{y}_i)\}_{i=1}^N$ - training data and labeled ground truth. Initialize $\boldsymbol{\theta}$ of feature extractor with the pretrained weights in ImageNet. Randomly initialize $\boldsymbol{\alpha}$ and $\boldsymbol{\beta}$.

Output: Anomaly detection network φ , and Confidence prediction network ζ .

- 1: – **STEP1: Training anomaly detection network**
- 2: **while** not converge **do**
- 3: Randomly sample a batch with m samples, including a half positive cases a half negative cases
- 4: Randomly sample l normal cases to compute μ_R, σ_R of l anomaly scores
- 5: Compute the anomaly score ν_i via Eq. 1 for each sample \mathbf{x}_i
- 6: Compute the anomaly detection loss $\frac{1}{m} \sum_{i=1}^m \mathcal{L}_{ano}(\nu_i, y_i, \mu_R, \sigma_R)$
- 7: Update parameters $\boldsymbol{\theta}$ and $\boldsymbol{\alpha}$ by using back-propagation algorithm
- 8: **end while**
- 9: – **STEP2: Training confidence prediction network**
- 10: Fix the feature extractor $\boldsymbol{\theta}$ and anomaly detection module $\boldsymbol{\alpha}$
- 11: **while** not converge **do**
- 12: Randomly sample a batch with m samples
- 13: Compute the confidence score ι_i via Eq. 6 for each sample \mathbf{x}_i
- 14: Generate the anomaly probability g_i for each sample
- 15: Compute the confidence prediction loss $\frac{1}{m} \sum_{i=1}^m \mathcal{L}_{conf}(\iota_i, g_i)$
- 16: Update parameters $\boldsymbol{\beta}$ by using back-propagation algorithm.
- 17: **end while**
- 18: – **STEP3: Joint training in an end-to-end manner**
- 19: **while** not converge **do**
- 20: Randomly sample a batch with m samples
- 21: Compute μ_R, σ_R
- 22: Compute ν_i and ι_i , then generate g_i according to Eq. 5
- 23: Update parameters $\boldsymbol{\theta}, \boldsymbol{\alpha}$, and $\boldsymbol{\beta}$ by using back-propagation in an end-to-end manner.
- 24: **end while**

anomaly detection network, which is the combination of the anomaly detection module and shared feature extractor, via minimizing the loss \mathcal{L}_{ano} using the vanilla stochastic gradient descent algorithm. Second, we fix feature extractor and train the confidence prediction network, which is the combination of the confidence prediction module and shared feature extractor, via minimizing the loss \mathcal{L}_{conf} using the Adam algorithm. In these two steps, we set the batch size to 40 and set the initial learning rate to 5×10^{-4} with linearly decay to 10^{-6} for the entire the training process. Third, we fine-tune the entire model in an end-to-end manner via minimizing the combination of \mathcal{L}_{ano} and \mathcal{L}_{conf} with the Adam optimizer. In this step, we set the batch size to 40 and set the initial

TABLE II
COMPARISON OF RESULTS OF BINARY CLASSIFICATION MODELS AND THE ANOMALY DETECTION MODEL ON THE X-VIRAL DATASET.

Mode		Feature extractor	Performance %			
Binary classification	Anomaly detection		Accuracy	Sensitivity	Specificity	AUC
✓		ResNet	78.52	78.28	78.56	86.24
	✓	ResNet	80.04	84.44	79.34	87.18
✓		EfficientNet	78.71	79.09	78.65	86.30
	✓	EfficientNet	80.65	85.51	79.87	87.42

learning rate to 10^{-7} , aiming to avoid deviating too much from the original anomaly detection scores. Note that the shared feature extractor are pre-trained on ImageNet, and the anomaly detection module and confidence prediction module are randomly initialized. The details of this training process is summarized in Algorithm 1.

In the inference stage, we input a test image into the well-trained model and generate a scalar anomaly score ν and a confidence score ι via the forward propagation. For the anomaly detection, we assume $g = 0.5$ in Eq. 5 as the boundary point. According to Eq. 4 and Eq. 5, we can compute the corresponding boundary anomaly score $\nu \approx 1.18$. Therefore, we set $T_{ano} = 1.18$ as the threshold to distinguish the abnormal cases and normal cases, i.e., detected as abnormal if $\nu \geq T_{ano}$ or otherwise. As for the confidence prediction, we empirically set $T_{conf} = 0.9$ as the threshold to correct the erroneous predictions with low confidence. In practice, we only re-label the predictions, recognized as normal or abnormal, with low confidence, i.e., $\iota < T_{conf}$, as abnormal cases to achieve the high sensitivity, which is significant in clinical study. Therefore, the final diagnosis made by our CAAD model is formulated as:

$$diag = \begin{cases} 1 & \text{if } \nu \geq T_{ano} \text{ or } \iota < T_{conf} \\ 0 & \text{if } \nu < T_{ano} \text{ and } \iota \geq T_{conf} \end{cases}. \quad (8)$$

Specifically, if either condition is met (i.e., the anomaly score is larger T_{ano} or the confidence score is less than T_{conf}), our model gives a 1 : POSITIVE diagnosis and recommend to be further examed by radiologists; otherwise our model gives a 0 : NEGATIVE diagnosis.

E. Performance metrics

For this study, the diagnostic performance of an algorithm is quantitatively assessed by the area under the receiver operator curve (AUC), sensitivity, specificity, and, accuracy. AUC reflects the probability that a recognition model ranks a randomly chosen positive instance higher than a randomly chosen negative case. It is the most commonly used metric to evaluate the overall classification performance. The sensitivity and specificity give the proportion of positives and negatives that were correctly identified, respectively. Accuracy gives the percentage of correctly classified cases, including both positives and negatives ones.

IV. DATASETS

Two in-house X-ray image datasets, X-VIRAL and X-COVID, were used for this study. The X-VIRAL dataset contains 5,977 viral pneumonia cases, 18,619 non-viral pneumonia cases, and 18,774 healthy controls (i.e., 5977 positive and

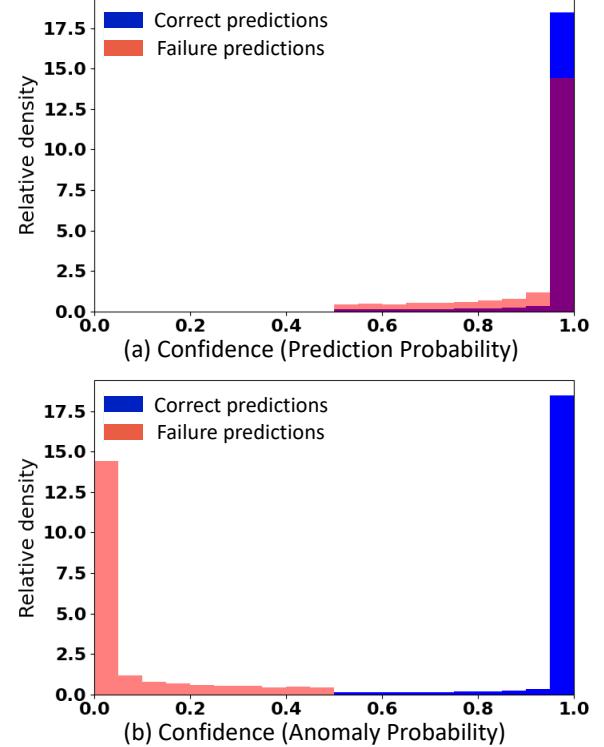


Fig. 3. Comparison of confidence learning based on (a) the prediction probability and (b) anomaly probability for failure prediction on the validation set of X-VIRAL.

37,393 negative cases) collected from 390 township hospitals through a telemedicine platform of JF Healthcare during 2019. Each X-ray image has a high resolution, varying from 1000 to 3000 for height and weight, and was annotated by one of three board certified radiologists. Note that all viral pneumonia cases were collected before the COVID-19 outbreak, and hence do not contain any COVID-19 cases. The X-COVID dataset was collected from 6 institutions during March 2020. It consists of 106 confirmed COVID-19 cases and 107 normal controls.

V. EXPERIMENTS AND RESULTS

A. Anomaly detection vs. binary classification

We first compared anomaly detection models to binary classification models on the X-VIRAL dataset using the five-fold cross validation. The feature extractor used in these models is either the 18-layer ResNet or EfficientNet-B0, both being pre-trained on ImageNet. The obtained performance metrics were listed in Table II. It shows that the binary classification using ResNet achieves the baseline performance, i.e., an

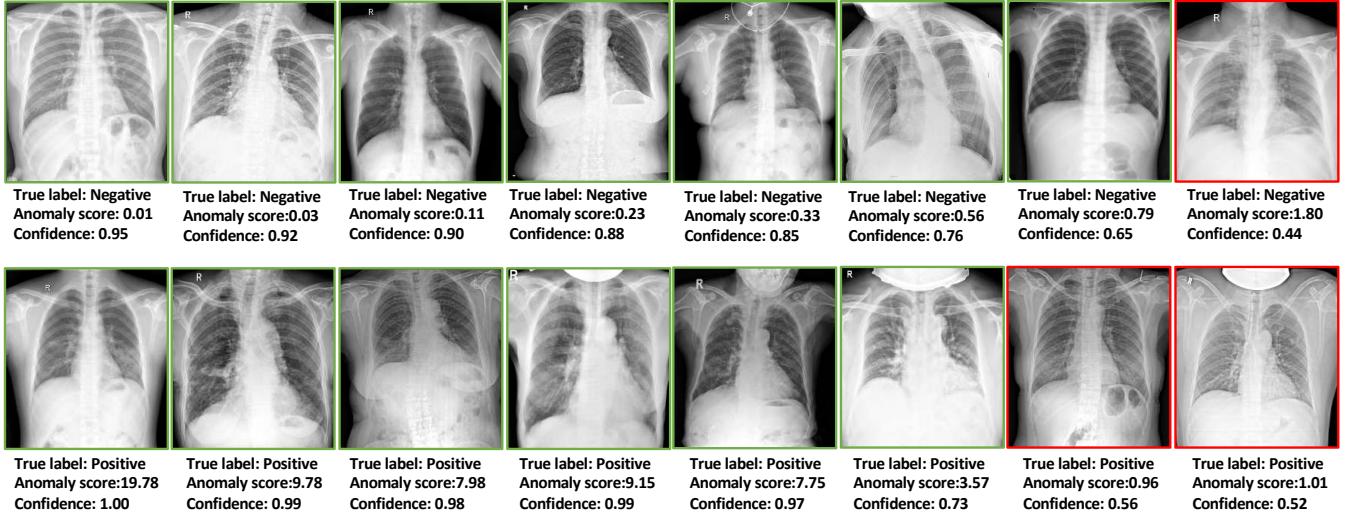


Fig. 4. A set of 16 chest x-ray images from the X-VIRAL validation set. The predicted anomaly score and confidence score are displayed beneath each image. Positive cases are shown in the top row, and negative cases are shown in the bottom row. Totally, 13 cases (marked with green boxes) were corrected diagnosed by our CAAD model, and three (marked with red boxes) were wrongly diagnosed.

accuracy of 78.52%, a sensitivity of 78.28%, a specificity of 78.56%, and an AUC of 86.24, which can be improved a little bit by introducing a stronger feature extractor, *i.e.*, EfficientNet-B0. By contrast, an anomaly detection model always outperforms (particularly in terms of sensitivity) the corresponding binary classification model. In this experiment, the anomaly detection model using EfficientNet-B0 achieves the highest accuracy of 80.65%, highest specificity of 79.87%, and highest AUC of 87.42%, and also remarkably improves the sensitivity from 79.09% to 85.51%. The results suggest that anomaly detection has distinct advantages over binary classification in distinguishing viral pneumonia cases from non-viral pneumonia cases and healthy controls, especially with an extraordinary ability to detect positive cases as evidenced by a high sensitivity, which is particularly useful for viral pneumonia screening in clinical practice.

B. Confidence learning for failure prediction

To analyze the effectiveness of confidence learning using either the anomaly probability or prediction probability, we compared the distribution of prediction probability and anomaly probability obtained on the validation set of X-VIRAL in Fig. 3. It is apparent that the prediction probability distributions of correct predictions and false predictions have a lot of overlap, which hinders the confidence prediction module from accurately distinguishing them. By contrast, the proposed anomaly probability can successfully separate false predictions from correct ones.

In Fig. 3(b) and Eq. 5, we observed that, if (1) the predicted anomaly score for negative cases is close to 0 or (2) the predicted anomaly score for positive cases is larger than T_{ano} , our CAAD model has a high confidence; otherwise low confidence. Such conjecture was confirmed with the results given in Fig. 4, which shows 16 chest X-ray images from the X-VIRAL validation set, each being equipped with the

TABLE III
PERFORMANCE OF THE ANODET MODEL (WITHOUT CONFIDENCE PREDICTION) AND OUR CAAD MODEL (WITH VARIABLE CONFIDENCE THRESHOLD T_{conf}) ON THE X-VIRAL DATASET.

Methods	T_{conf}	Accuracy	Sensitivity	Specificity	AUC
AnoDet	/	80.65	85.51	79.87	87.42
	0.5	80.33	85.88	79.44	87.57
	0.6	79.47	87.02	78.27	
	0.7	78.48	88.05	76.95	
	0.8	76.79	89.63	74.74	
	0.9	71.21	93.01	67.72	
	0.95	46.44	97.69	38.25	

ground truth label and the anomaly scores and confidence scores predicted by our CAAD model. For each negative case in the top row, it shows that if the predicted anomaly scores are very low, close to 0, the confidence score is close to 1. Similarly, for each positive case in the bottom row, it shows that if the predicted anomaly score is larger than T_{conf} , the confidence score is still close to 1. In contrast, the confidence becomes very low if a case is wrongly diagnosed, as those marked with red bounding boxes. These results demonstrate the effectiveness of using the proposed anomaly probability to learn the confidence for failure prediction.

C. Importance of confidence prediction

To evaluate the effectiveness of confidence prediction, we compared the anomaly detection network (denoted by AnoDet) with our CAAD model, in which the confidence threshold T_{conf} ranges from 0.5 to 0.95. Note that the CAAD model can be treated as the combination of AnoDet with a confidence prediction module. The results obtained on the X-VIRAL dataset are shown in Table III. It reveals that (1) using confidence prediction leads to a slightly improved AUC of 87.57%, improved sensitivity, and deteriorated accuracy and specificity; (2) when setting the confidence threshold

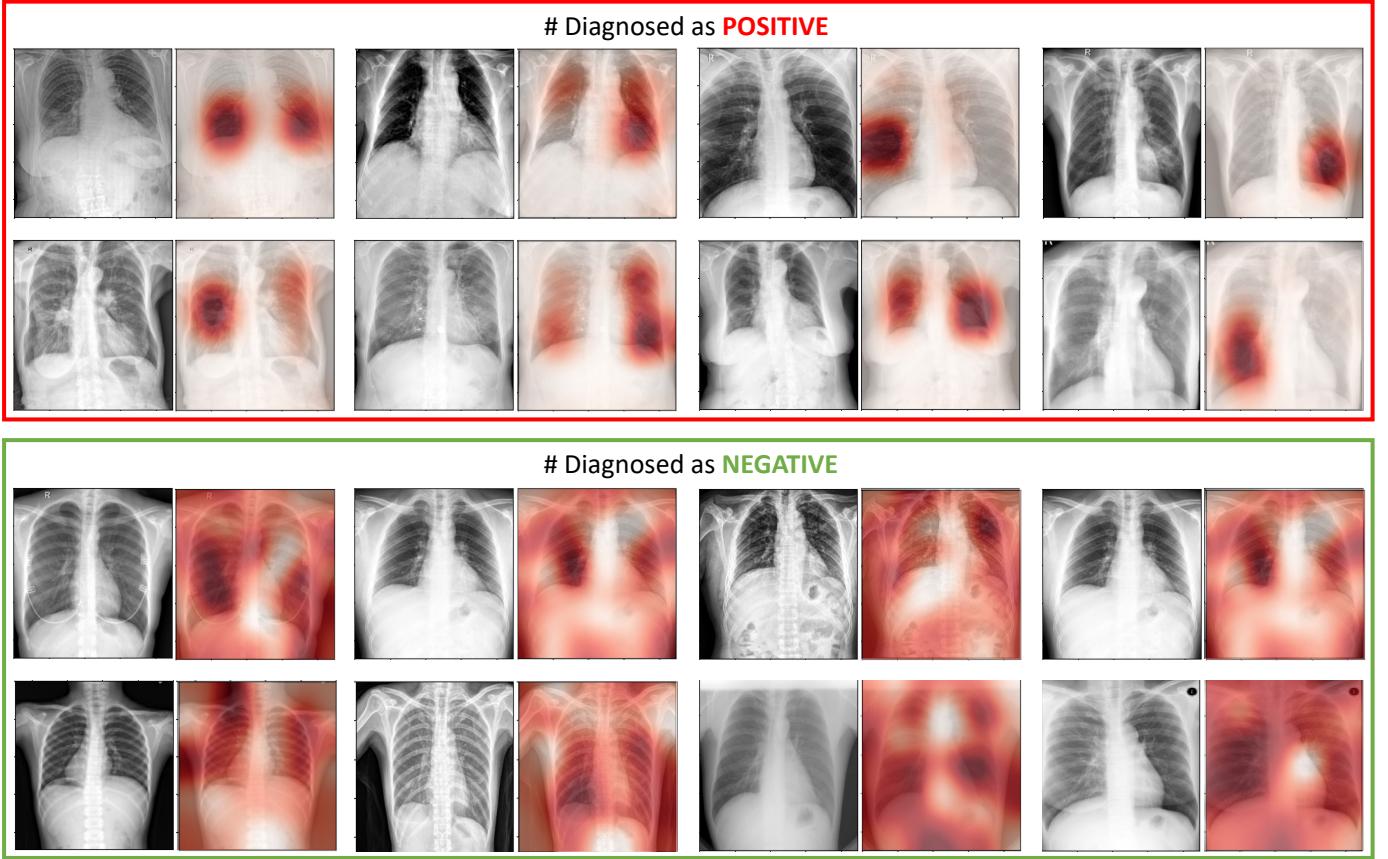


Fig. 5. Visualization of 16 chest X-ray images from the X-VIRAL validation set and their Grad-CAMs. The images in two top rows are diagnosed as positive by our CAAD model, while the images in two bottom rows are diagnosed as negative.

T_{conf} to 0.5, the performance of our model is comparable to that of AnoDet; and (3) with the increase of the confidence threshold T_{conf} , the deterioration of accuracy and specificity becomes severe and the improvement of sensitivity becomes substantial. Specifically, when setting the confidence threshold T_{conf} to 0.9, our CAAD model can boost the sensitivity from 85.51% to 93.01% while having a specificity of 67.72%. The improvement in the sensitivity suggests the advantage of using confidence prediction in our model for the clinical screening of viral pneumonia. It should be noted that, as mentioned in Section III-C, we propose the confidence prediction module to distinguish the successful predictions with a high confidence from failed predictions with a low confidence, instead of further improving the performance of anomaly detection.

D. Visualizing region of diagnosis

For the visual explanation of the decision reasoning of our CAAD model, we adopted the gradient-weighted class activation mapping (Grad-CAM) [32] to "see" which regions play an important role during the inference. Fig. 5 shows 16 chest X-ray images from the X-VIRAL validation set, each being accompanied with the Grad-CAM maps overlaid on it. Eight cases in two top rows were diagnosed as positive by our CAAD model, while the other cases in two bottom rows were diagnosed as negative. It reveals that our CAAD model is able to focus on the suspected lesions and, accordingly,

TABLE IV
COMPARING PERFORMANCE OF FOUR MODELS ON THE X-COVID
DATASET.

Model	Performance %			
	Accuracy	Sensitivity	Specificity	AUC
Efficientnet [25]	69.95	45.28	94.39	74.45
ConfidNet [23]	68.08	69.81	66.36	74.89
AnoDet [17]	73.24	55.66	90.65	82.97
CAAD	72.77	71.70	73.83	83.61

diagnose the input image as POSITIVE. However, if no highly suspected lesion is found, our model assigns the homogeneous activation values to almost the entire image. The visualization demonstrates the good interpretability of our CAAD model.

E. Model generalization

To demonstrate its generalization ability, the well-trained CAAD model was directly tested on the unseen X-COVID dataset without fine-tuning. It was also compared to a binary classifier using EfficientNet [25], a binary classifier with confidence prediction (*i.e.*, ConfidNet) [23], and a anomaly detection model (*i.e.*, AnoDet) [17]. Note that all of these models were never trained on the COVID-19 cases. Table. IV gives the performance of these models on the X-COVID dataset. It reveals that (1) anomaly detection models are superior to both binary classifiers, especially in terms of AUC,

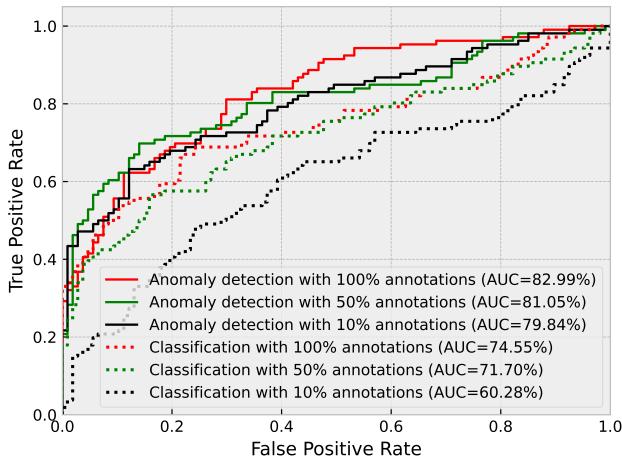


Fig. 6. AUC curves of anomaly detection and binary classification models obtained on X-COVID when both models were trained with 100%, 50% and 10% annotations on X-VIRAL.

which demonstrates the effectiveness of anomaly detection in viral pneumonia screening; (2) confidence prediction is an effective strategy to predict failures in both a classifier and an anomaly detector, contributing to a big improvement in sensitivity; and (3) the performance of these models, however, drops when comparing to their performance in viral pneumonia screening. In summary, we expect that our model trained on the X-VIRAL dataset would have the ability to detect unseen COVID-19 cases as anomaly. Our results suggests that, despite the significant drop of sensitivity, our CAAD model still achieves an AUC of 83.61% and a sensitivity of 71.70% for COVID-19 screening, the highest performance obtained in our experiment.

F. Learning with less positive samples

For real applications, it is much more difficult to collect positive (*i.e.*, viral pneumonia) samples than to collect negative samples. Hence automated viral pneumonia screening is a class-imbalance problem. Anomaly detection methods can better handle the class-imbalance issue than classification methods. To verify this, we trained the anomaly detection model and binary classification model, both using EfficientNet as the feature extractor, under more imbalanced conditions, *i.e.*, learning on less positive samples, and then tested them without any fine-tuning on the X-COVID dataset. Specifically, in the training stage, we kept the number of negative samples unchanged and reduced the positive samples from 100% to 50% and 10%, respectively. The obtained AUC curves were shown in Fig. 6. It revels that training with less positive data leads to the poor performance for both anomaly detection and binary classification models. However, anomaly detection has the less performance degradation than binary classification, which indicates the superiority of anomaly detection in terms of addressing imbalanced problem.

VI. DISCUSSION

In the escalation of the COVID-19 epidemic, many attempts have been made to develop fast and accurate COVID-19 screening by means of chest medical imaging [33], [34], [35], [36], [37]. Kang *et al.* [36] leveraged different types of features extracted from chest CT images, and introduced a multi-view representation learning method to distinguish positive COVID-19 from negative cases. Shi *et al.* [37] presented an infection size aware random forest model to classify 1658 positive subjects confirmed COVID-19 and 1027 negative subjects without COVID-19 infection using chest CT, and achieved the sensitivity of 90.70% and specificity of 83.30% on this large-scale CT dataset. However, CT imaging takes considerably more time than X-ray imaging, and needs more complex sanitization procedures between switching patients. Besides, sufficient high-quality CT scanners may not be available in many under-developed regions, making it difficult for a timely viral pneumonia screening. In contrast, X-ray imaging is the most common and widely available chest imaging technique, playing a crucial role in clinical care and epidemiological studies [38], [39]. Most ambulatory care facilities, even in rural regions, has X-ray imaging capability. Besides, X-ray imaging is real-time which can significantly speed up the screening of a mass population in a relatively short time and at a significantly reduced cost. Wang *et al.* [40] and Apostolopoulos *et al.* [41] introduced DCNN-based binary classification models for the detection of COVID-19 cases using chest X-Ray imaging.

Different from these COVID-19 screening works, we view COVID-19 as a novel type of viral pneumonia and attempt to distinguish it, together with other types of viral pneumonia, from non-viral pneumonia and healthy controls. To this end, we reformulate the binary classification problem in an anomaly detection fashion. Besides, we introduce a confidence prediction module to estimate the reliability of model diagnosis by learning an anomaly probability as the model confidence. The proposed CAAD model achieves an AUC of 83.61% on COVID-19 screening, which outperforms other AI based methods [42]. Although achieving a sensitivity of only 71.70%, our CAAD model shows a screening ability that is comparable to that of radiologists, as a sensitivity of 69% was reported in [18]. The reason of such a low sensitivity may attribute to the observation that some subjects may have not developed radiographic visible pathology in their lungs at the early stage of viral pneumonia when the X-ray was taken [43].

VII. CONCLUSION

In this paper, we propose the CAAD model for viral pneumonia screening. Our results on two chest X-ray datasets indicate that (1) anomaly detection works well in term of viral pneumonia screening on chest X-ray images and is superior to binary classification methods, and (2) learning model confidence is useful to predict failures, greatly reducing the false negatives, and (3) our CAAD model, never seeing any COVID-19 cases, achieves an AUC of 83.61% and a sensitivity of 71.70% on the unseen X-COVID dataset, which is comparable to the performance of medical professionals.

Our future work will focus on further reducing the false negative rate and, if possible, decreasing the false positive rate as well. We will also investigate how to differentiate the viral pneumonia severity using chest X-ray and then detect the potentially severe cases for early interventions, which requires more clinical diagnostic information.

REFERENCES

- [1] Y. Bai, L. Yao, T. Wei, F. Tian, D.-Y. Jin, L. Chen, and M. Wang, "Presumed asymptomatic carrier transmission of COVID-19," *Journal of the American Medical Association (JAMA)*, 2020.
- [2] H. Chen, J. Guo, C. Wang, F. Luo, X. Yu, W. Zhang, J. Li, D. Zhao, D. Xu, Q. Gong *et al.*, "Clinical characteristics and intrauterine vertical transmission potential of covid-19 infection in nine pregnant women: a retrospective review of medical records," *The Lancet*, vol. 395, no. 10226, pp. 809–815, 2020.
- [3] W. Li, M. J. Moore, N. Vasilieva, J. Sui, S. K. Wong, M. A. Berne, M. Somasundaran, J. L. Sullivan, K. Luzuriaga, T. C. Greenough *et al.*, "Angiotensin-converting enzyme 2 is a functional receptor for the sars coronavirus," *Nature*, vol. 426, no. 6965, pp. 450–454, 2003.
- [4] E. I. Azhar, S. A. El-Kafrawy, S. A. Farraj, A. M. Hassan, M. S. Al-Saeed, A. M. Hashem, and T. A. Madani, "Evidence for camel-to-human transmission of mers coronavirus," *New England Journal of Medicine*, vol. 370, no. 26, pp. 2499–2505, 2014.
- [5] V. M. Corman, O. Landt, M. Kaiser, R. Molenkamp, A. Meijer, D. K. Chu, T. Bleicker, S. Brünink, J. Schneider, M. L. Schmidt *et al.*, "Detection of 2019 novel coronavirus (2019-ncov) by real-time rt-pcr," *Eurosurveillance*, vol. 25, no. 3, p. 2000045, 2020.
- [6] L. Lan, D. Xu, G. Ye, C. Xia, S. Wang, Y. Li, and H. Xu, "Positive rt-pcr test results in patients recovered from covid-19," *Jama*, vol. 323, no. 15, pp. 1502–1503, 2020.
- [7] W. H. Self, D. M. Courtney, C. D. McNaughton, R. G. Wunderink, and J. A. Kline, "High discordance of chest x-ray and computed tomography for detection of pulmonary opacities in ed patients: implications for diagnosing pneumonia," *The American journal of emergency medicine*, vol. 31, no. 2, pp. 401–405, 2013.
- [8] R. Joarder and N. Crundwell, *Chest X-ray in clinical practice*. Springer Science & Business Media, 2009.
- [9] X. Wang, Y. Peng, L. Lu, Z. Lu, M. Bagheri, and R. M. Summers, "Chestx-ray8: Hospital-scale chest x-ray database and benchmarks on weakly-supervised classification and localization of common thorax diseases," in *Proceedings of IEEE Conference on Computer Vision and Pattern Recognition*, 2017, pp. 2097–2106.
- [10] P. Rajpurkar, J. Irvin, K. Zhu, B. Yang, H. Mehta, T. Duan, D. Ding, A. Bagul, C. Langlotz, K. Shpanskaya *et al.*, "Chexnet: Radiologist-level pneumonia detection on chest x-rays with deep learning," *arXiv preprint arXiv:1711.05225*, 2017.
- [11] H. Wang, H. Jia, L. Lu, and Y. Xia, "Thorax-net: An attention regularized deep neural network for classification of thoracic diseases on chest radiography," *IEEE journal of biomedical and health informatics*, 2019.
- [12] D. S. Kermany, M. Goldbaum, W. Cai, C. C. Valentim, H. Liang, S. L. Baxter, A. McKeown, G. Yang, X. Wu, F. Yan *et al.*, "Identifying medical diagnoses and treatable diseases by image-based deep learning," *Cell*, vol. 172, no. 5, pp. 1122–1131, 2018.
- [13] S. Rajaraman, S. Candemir, I. Kim, G. Thoma, and S. Antani, "Visualization and interpretation of convolutional neural network predictions in detecting pneumonia in pediatric chest radiographs," *Applied Sciences*, vol. 8, no. 10, p. 1715, 2018.
- [14] C. Bellinger, S. Sharma, and N. Japkowicz, "One-class versus binary classification: Which and when?" in *2012 11th International Conference on Machine Learning and Applications*, vol. 2. IEEE, 2012, pp. 102–106.
- [15] C. Zhou and R. C. Paffenroth, "Anomaly detection with robust deep autoencoders," in *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2017, pp. 665–674.
- [16] G. Pang, L. Cao, L. Chen, and H. Liu, "Learning representations of ultrahigh-dimensional data for random distance-based outlier detection," in *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2018, pp. 2041–2050.
- [17] G. Pang, C. Shen, and A. van den Hengel, "Deep anomaly detection with deviation networks," in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2019, pp. 353–362.
- [18] H. Y. F. Wong, H. Y. S. Lam, A. H.-T. Fong, S. T. Leung, T. W.-Y. Chin, C. S. Y. Lo, M. M.-S. Lui, J. C. Y. Lee, K. W.-H. Chiu, T. Chung *et al.*, "Frequency and distribution of chest radiographic findings in covid-19 positive patients," *Radiology*, p. 201160, 2020.
- [19] H. Magree, F. Russell, R. Sa'aga, P. Greenwood, L. Tikoduadua, J. Pryor, L. Waqatakirewa, J. Carapetis, and E. K. Mulholland, "Chest x-ray-confirmed pneumonia in children in fiji," *Bulletin of the World Health Organization*, vol. 83, pp. 427–433, 2005.
- [20] A. Jacobi, M. Chung, A. Bernheim, and C. Eber, "Portable chest x-ray in coronavirus disease-19 (covid-19): A pictorial review," *Clinical Imaging*, 2020.
- [21] A. Borghesi and R. Maroldi, "Covid-19 outbreak in italy: experimental chest x-ray scoring system for quantifying and monitoring disease progression," *La radiologia medica*, p. 1, 2020.
- [22] D. Hendrycks and K. Gimpel, "A baseline for detecting misclassified and out-of-distribution examples in neural networks," in *International Conference on Learning Representation*, 2017.
- [23] C. Corbière, N. Thome, A. Bar-Hen, M. Cord, and P. Pérez, "Addressing failure prediction by learning model confidence," in *Advances in Neural Information Processing Systems*, 2019, pp. 2898–2909.
- [24] Y. Xie, H. Lu, J. Zhang, C. Shen, and Y. Xia, "Deep segmentation-emendation model for gland instance segmentation," in *International Conference on Medical Image Computing and Computer-Assisted Intervention*. Springer, 2019, pp. 469–477.
- [25] M. Tan and Q. Le, "Efficientnet: Rethinking model scaling for convolutional neural networks," in *International Conference on Machine Learning*, 2019, pp. 6105–6114.
- [26] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, "Imagenet: A large-scale hierarchical image database," in *Proceedings of IEEE Conference on Computer Vision and Pattern Recognition*. Ieee, 2009, pp. 248–255.
- [27] M. Sandler, A. Howard, M. Zhu, A. Zhmoginov, and L.-C. Chen, "Mobilenetv2: Inverted residuals and linear bottlenecks," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 4510–4520.
- [28] M. Tan, B. Chen, R. Pang, V. Vasudevan, M. Sandler, A. Howard, and Q. V. Le, "Mnasnet: Platform-aware neural architecture search for mobile," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2019, pp. 2820–2828.
- [29] J. Hu, L. Shen, and G. Sun, "Squeeze-and-excitation networks," in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2018, pp. 7132–7141.
- [30] R. Hadsell, S. Chopra, and Y. LeCun, "Dimensionality reduction by learning an invariant mapping," in *Proceedings of IEEE Conference on Computer Vision and Pattern Recognition*, vol. 2. IEEE, 2006, pp. 1735–1742.
- [31] T. Chen, S. Kornblith, M. Norouzi, and G. Hinton, "A simple framework for contrastive learning of visual representations," *arXiv:2002.05709*, 2020.
- [32] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, "Grad-cam: Visual explanations from deep networks via gradient-based localization," in *Proceedings of the IEEE international conference on computer vision*, 2017, pp. 618–626.
- [33] D. S. W. Ting, L. Carin, V. Dzau, and T. Y. Wong, "Digital technology and covid-19," *Nature medicine*, vol. 26, no. 4, pp. 459–461, 2020.
- [34] L. Huang, R. Han, T. Ai, P. Yu, H. Kang, Q. Tao, and L. Xia, "Serial quantitative chest ct assessment of covid-19: Deep-learning approach," *Radiology: Cardiothoracic Imaging*, vol. 2, no. 2, p. e200075, 2020.
- [35] S. Wang, B. Kang, J. Ma, X. Zeng, M. Xiao, J. Guo, M. Cai, J. Yang, Y. Li, X. Meng *et al.*, "A deep learning algorithm using ct images to screen for corona virus disease (covid-19)," *MedRxiv*, 2020.
- [36] H. Kang, L. Xia, F. Yan, Z. Wan, F. Shi, H. Yuan, H. Jiang, D. Wu, H. Sui, C. Zhang *et al.*, "Diagnosis of coronavirus disease 2019 (covid-19) with structured latent multi-view representation learning," *IEEE transactions on medical imaging*, 2020.
- [37] F. Shi, L. Xia, F. Shan, D. Wu, Y. Wei, H. Yuan, H. Jiang, Y. Gao, H. Sui, and D. Shen, "Large-scale screening of covid-19 from community acquired pneumonia using infection size-aware classification," 2020.
- [38] T. Cherian, E. K. Mulholland, J. B. Carlin, H. Ostensen, R. Amin, M. d. Campo, D. Greenberg, R. Lagos, M. Lucero, S. A. Madhi *et al.*, "Standardized interpretation of paediatric chest radiographs for the diagnosis of pneumonia in epidemiological studies," *Bulletin of the World Health Organization*, vol. 83, pp. 353–359, 2005.
- [39] T. Franquet, "Imaging of pneumonia: trends and algorithms," *European Respiratory Journal*, vol. 18, no. 1, pp. 196–208, 2001.

- [40] L. Wang and A. Wong, "Covid-net: A tailored deep convolutional neural network design for detection of covid-19 cases from chest radiography images," *arXiv*, pp. arXiv-2003, 2020.
- [41] I. D. Apostopoulos and T. A. Mpesiana, "Covid-19: automatic detection from x-ray images utilizing transfer learning with convolutional neural networks," *Physical and Engineering Sciences in Medicine*, p. 1, 2020.
- [42] K. Murphy, H. Smits, A. J. Knoops, M. B. Korst, T. Samson, E. T. Scholten, S. Schalekamp, C. M. Schaefer-Prokop, R. H. Philipsen, A. Meijers *et al.*, "Covid-19 on the chest radiograph: A multi-reader evaluation of an ai system," *Radiology*, p. 201874, 2020.
- [43] M. B. Weinstock, A. Echenique, J. W. R. DABR, A. Leib, and F. A. ILLUZZI, "Chest x-ray findings in 636 ambulatory patients with covid-19 presenting to an urgent care center: a normal chest x-ray is no guarantee," *J Urgent Care Med*, vol. 14, no. 7, pp. 13–8, 2020.